



UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Distributed Computing - MPI

Dimitar Lukarski, Jarmo Rantakokko

Division of Scientific Computing
Department of Information Technology
Uppsala Programming for Multicore Architectures Research Center
(UPMARC)
Uppsala University

Programming of Parallel Computers, Jan, 2014

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

What You Should Know...

- ▶ Programming in FORTRAN/C/C++/Java
- ▶ Basics in hardware - CPU, RAM, Network

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Outline

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala

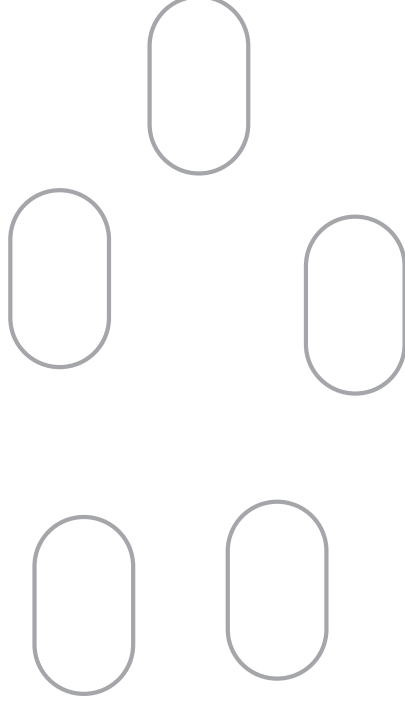


UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Distributed computing



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

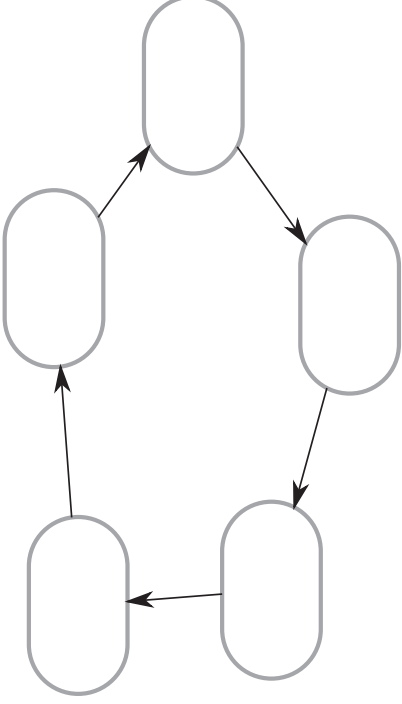
Global Funcs

Timing

Datatypes

Topology

Distributed computing



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

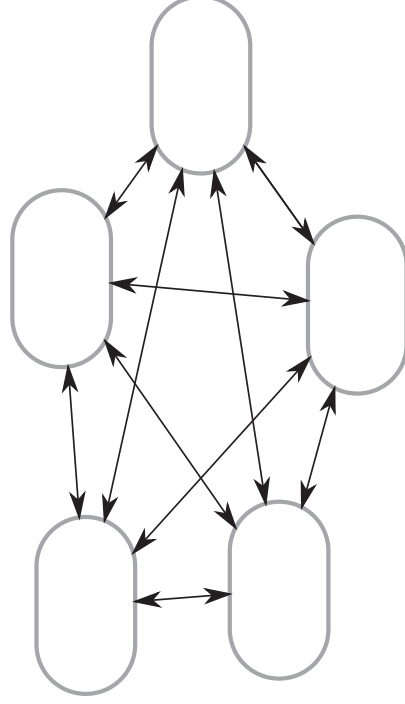
Global Funcs

Timing

Datatypes

Topology

Distributed computing



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





- Outline
- Intro and motivation**
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology



- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology

- ▶ Sending data
- ▶ Receiving data
- ▶ Waiting for data
- ▶ Waiting for synchronization

D. Lukarski, J. Rantakokko, Jan, 2014, Uppsala



- ▶ MPI is a library for data and messages passing
- ▶ The communication is explicit (the user specify what/how needs to be sent)
- ▶ The number of processes is configured at start up

MPI program:

- ▶ Communication functions (Send/Receive/etc)
- ▶ C, C++, FORTRAN (+ interface to others)
- ▶ To run the program you need a run-time agent (mpirun)
- ▶ Example `mpirun -np 2 ./program`

- ▶ Communication functions (Send/Receive/etc)

- ▶ C, C++, FORTRAN (+ interface to others)
- ▶ To run the program you need a run-time agent (mpirun)
- ▶ Example `mpirun -np 2 ./program`



- ▶ 1992 - draft of the project
- ▶ 1994 - first version MPI 1.0
 - ▶ C, FORTRAN 77
 - ▶ Point-to-point
 - ▶ Global communication, groups
- ▶ 1997 - MPI 2.0
 - ▶ One-sided communication
 - ▶ C++, FORTRAN 90
 - ▶ Dynamic management
- ▶ 2008/09 - MPI 2.1 and 2.2
- ▶ 2012 - MPI-3.0
 - ▶ New one-sided communication

► One-sided communication

- ▶ C++, FORTRAN 90
- ▶ Dynamic management

- 2008/09 - MPI 2.1 and 2.2

► 2012 - MPI-3.0

- ▶ New one-sided communication



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

The Body of Your Code / C, C++

The program should contain

- ▶ *#include <mpi.h>*
 - ▶ the include file for MPI
- ▶ `MPI_Init(&agrc, &argv);`
 - ▶ Initialize the MPI
 - ▶ Pass the arguments from the command line
- ▶ `MPI_Finalize();`
 - ▶ Finalize all communication

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Compile and Run Your Code / Linux

Compilation

- ▶ *mpicc -O3 -o program program.c*
- ▶ *mpiCC -O3 -o program program.cpp*

Execution

- ▶ *mpirun -np N ./program*
- ▶ where *N* is the number of process (integer number)

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





- Outline
- Intro and motivation
- Code Body
- Communicators**
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology

- ▶ Each process is associated with ID number - called *rank*
- ▶ Processes can belong to different groups

- ▶ The standard (default) communicator for all MPI process is called *MPI_COMM_WORLD*



- Outline
- Intro and motivation
- Code Body
- Communicators**
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology





UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

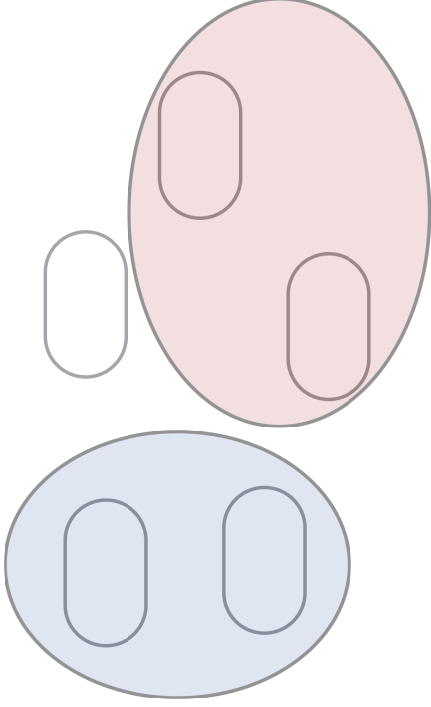
Global Funcs

Timing

Datatypes

Topology

Communicators



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

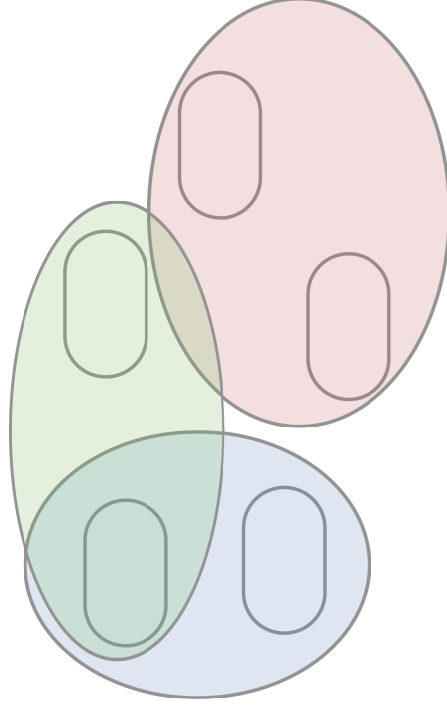
Global Funcs

Timing

Datatypes

Topology

Communicators



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Communicator Functions

Get the current rank (ID)

- ▶ *MPI_Comm_rank()*

Get the size of the communicator

- ▶ *MPI_Comm_size()*

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Code Example

```
#include <stdio.h>
#include <mpi.h>

int main (int argc,
          char *argv[]) {
    int rank, size;

    MPI_Init (&argc, &argv);
    MPI_Comm_rank (MPI_COMM_WORLD, &rank);
    MPI_Comm_size (MPI_COMM_WORLD, &size);

    printf("Hello world from process %d of %d\n",
          rank, size );

    MPI_Finalize();
    return 0;
}
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation
Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Send and Receive



- Point-to-point communication

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala

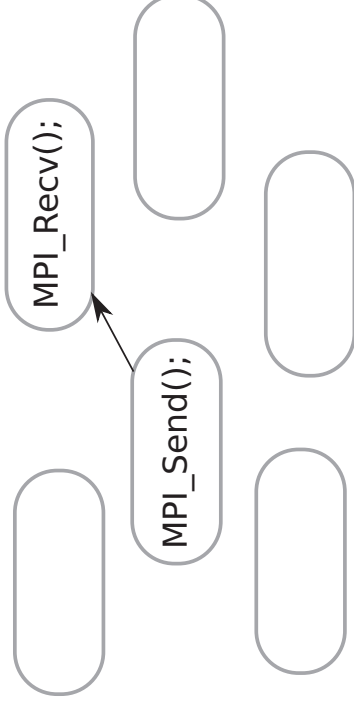




UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Send - Recv



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Send

```
MPI_Send( buff,
           count,
           datatype,
           dest,
           tag,
           comm);
```

void*	buff	Address of the buffer
int	count	Number of elements in the buff
MPI_Datatype	datatype	Data type(MPI_INT,...)
int	dest	Destination (rank)
int	tag	Tag
MPI_Comm	comm	MPI-Communicator

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Recv

```
MPI_Recv( buff,  
          count,  
          datatype,  
          source,  
          tag,  
          comm  
          status);
```

void*	buff	Address of the buffer
int	count	Number of elements in the buff
MPI_Datatype	datatype	Data type(MPI_INT,...)
int	src	Source (rank)
int	tag	Tag
MPI_Comm	comm	MPI-Communicator
MPI_status*	status	Status

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Code Example

The following code compute the following

- ▶ Proc 0 sends data to Proc 1
- ▶ Proc 1 sends data to Proc 2
- ▶ ...
- ▶ Proc N-1 sends data to Proc 0

The data is a double value, where each process adds its rank. Thus, at the end the data will contain the sum of all ranks $(0+1+2+3+4+...+N-1)$.

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





```
#include<stdio.h>
#include<mpi.h>
```

```
int main(int argc, char *argv[]) {
    int size, rank;
    MPI_Status status;
    double bot;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    bot = 0.0;
```



```

if (rank==0){
    MPI_Send(&bot,1,MPI_DOUBLE,1,1,MPI_COMM_WORLD);
    MPI_Recv(&bot,1,MPI_DOUBLE,size-1,1,
             MPI_COMM_WORLD,&status);
    printf("result: %lf\n", bot);
}

```

```

}else {
    MPI_Recv(&bot,1,MPI_DOUBLE,rank-1,1,
             MPI_COMM_WORLD,&status);
    bot += rank;
    MPI_Send(&bot,1,MPI_DOUBLE,
             (rank+1)%size,1,MPI_COMM_WORLD);
}
MPI_Finalize();
return 0;
}

```





(Non-)Blocking/(A-)Synchronous

Blocking mode

- ▶ The current process waits until the its part finished

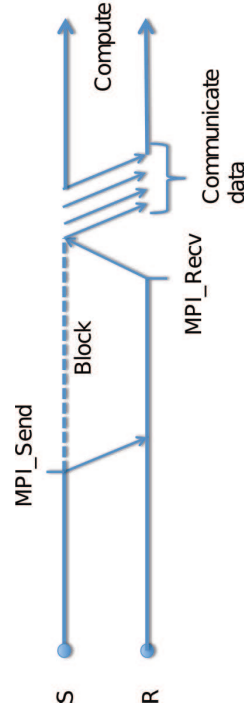
Synchronous mode

- ▶ The process waits for the start of the other side



MPI_Send

Blocking Send (Same as MPI_Ssend)

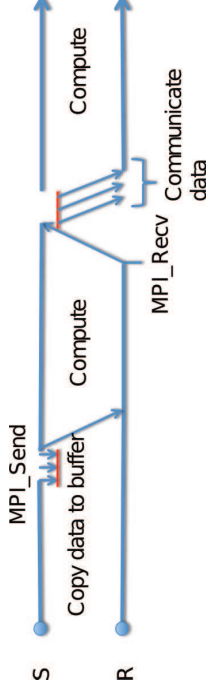


- ▶ The processor is blocked until data has been sent
- ▶ Receive can be posted before the send
- ▶ Receiver waits for the ready-to-send-signal
- ▶ Deadlock will occur if no matching send/receive



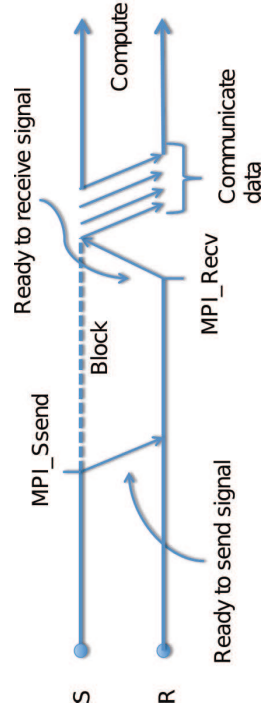
MPI_Send

When the data is very small (size < threshold)



MPI_Ssend

Blocking MPI_Ssend

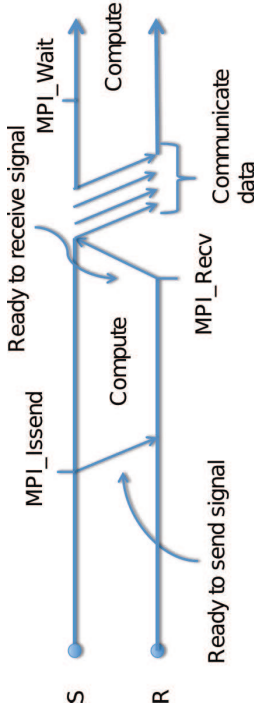


- It will always wait until the receive has been posted on the receiving end.



MPI_Issend

Non-blocking MPI_Issend

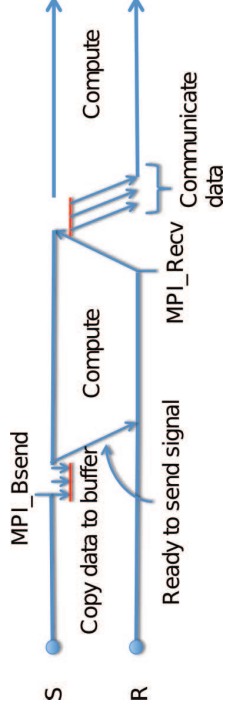


- ▶ The sender is not blocked but we need to do an explicit wait to ensure that the data buffer is safe to use (data has been transmitted) with MPI_Wait! Can also do a non-blocking MPI_Test.



MPI_Bsend

Buffered MPI_Bsend



- ▶ It copies data to a user-supplied communication buffer and then returns (no blocking of processor). It is safe to modify the original data and no need to wait for the communication to complete.

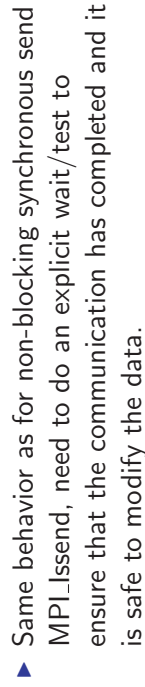


- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive**
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology

- ▶ `int buflen=totlen*sizeof(double) + MPI_BSEND_OVERHEAD;`
- ▶ `double *buffer=malloc(buflen);`
- ▶ `MPI_Buffer_attach(buffer,buflen);`
- ▶ `MPI_Bsend(data,count,type,dest,tag,comm);`



Non-blocking MPI_send

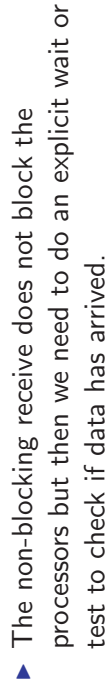




Ready send MPI_Rsend



Non-blocking MPI_recv





- ▶ MPI_Send, MPI_Isend – Standard mode
- ▶ MPI_Ssend, MPI_Issend – Synchronous mode
- ▶ MPI_Rsend, MPI_Irsend – Ready mode
- ▶ MPI_Bsend, MPI_Ibsend – Buffered mode

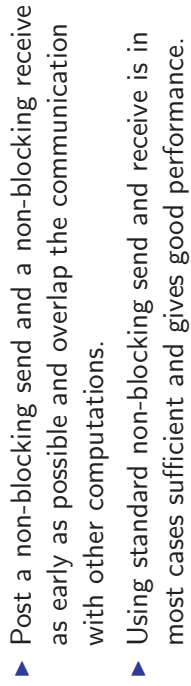
All modes can be blocked -or- non-blocked. Blocking suspends the execution until the message (data) buffer is safe to use (message has been sent/ received/copied). Non-blocking calls return immediately after initiating the communication.



All send calls need to be matched with a receive call!
and
Deadlock will occur if no matching send/receive!



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive**
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology



A vertical toolbar containing 15 icons for document editing and presentation navigation. From top to bottom, the icons are: a circular arrow (refresh), a magnifying glass (search), a curved arrow (undo), a list of four horizontal bars (table of contents), a triangle pointing up (previous slide), a list of four horizontal bars (table of contents), a triangle pointing down (next slide), a triangle pointing up (previous slide), a list of four horizontal bars (table of contents), a triangle pointing down (next slide), a stack of three squares (layers), a triangle pointing down (next slide), a triangle pointing up (previous slide), a square (insert), and a triangle pointing down (next slide).



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs**
- Global Funcs
- Timing
- Datatypes
- Topology



Other P2P Funcs

MPI_Wait

- MPI_Waitall

- MPI_Waitany

- MPI_Waitsome



Other P2P Funcs

MPI_Test, MPI_Testall, MPI_Testany, MPI_Testsome

- MPI_Probe, MPI_Iprobe

- MPI_Cancel

-



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Global Functions

Barrier

- ▶ All processes wait on it

Broadcast

- ▶ Send to all processes

Reduce

- ▶ Collect data with an operation

Scatter/Gather

- ▶ Send/Receive to/from all

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





- A global barrier

```
int MPI_Barrier(MPI_Comm comm);
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



- A processor sends data to all others





•
•
•

```
int global_var;
```

if (rank == 0)

```
global_var = read_file("var.dat");
```

MPI_Bcast(global_var, 1, MPI_INT, 0,

•
•
•



- Reduces values on all processes to a single value





```
...
// declare and compute a[]
...
double local_sum = 0.0;

for( int i = 0; i < n; i++)
    local_sum += a[i];

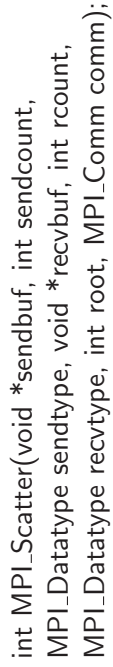
double total;
MPI_Reduce(sum, total, 1, MPI_DOUBLE, MPI_SUM,
            0, MPI_COMM_WORLD);
...
```



MPI_MAX	Maximum
MPI_MIN	Minimum
MPI_SUM	Sum
MPI_PROD	Product
MPI_LAND	Logical AND
MPI_BAND	Bitwise AND
MPI_LOR	Logical OR
MPI_BOR	Bitwise OR
MPI_LXOR	Logical XOR
MPI_BXOR	Bitwise XOR
MPI_MAXLOC	Max value and location
MPI_MINLOC	Min value and location



- ▶ A processor sends data to all others



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



```

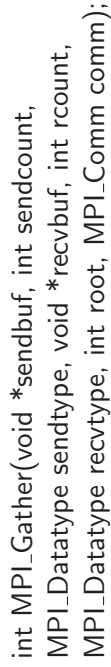
...
MPI_Comm comm;
int gsize,*sendbuf;
int root, rbuf[100];
...
MPI_Comm_rank( comm, myrank);
if ( myrank == root) {
    MPI_Comm_size( comm, &gsize);
    sendbuf = (int *)malloc(gsize*100*sizeof(int));
}
...
MPI_Scatter(sendbuf, 100, MPI_INT, rbuf, 100,
            MPI_INT, root, comm);
...

```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



- ▶ A processor sends data to all others



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



```

...
MPI_Comm comm;
int gsize, sendarray[100];
int root, myrank, *rbuf;

...
MPI_Comm_rank( comm, myrank);
if ( myrank == root) {
    MPI_Comm_size( comm, &gsize);
    rbuf = (int *)malloc(gsize*100*sizeof(int));
}

MPI_Gather(sendarray, 100, MPI_INT, rbuf, 100,
           MPI_INT, root, comm);

...

```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala

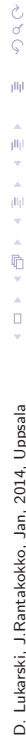




- ▶ MPI_Allreduce
- ▶ MPI_Allgather
- ▶ MPI_Alltoall



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs**
- Timing
- Datatypes
- Topology

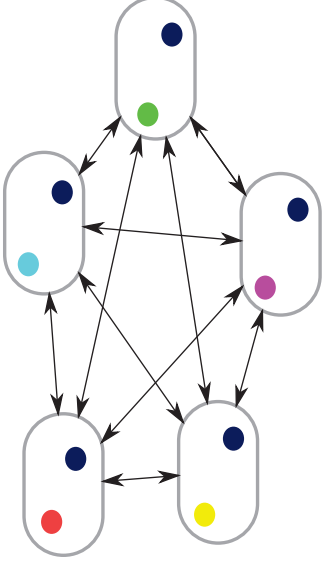




UPPSALA
UNIVERSITET

MPI_Allreduce

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology



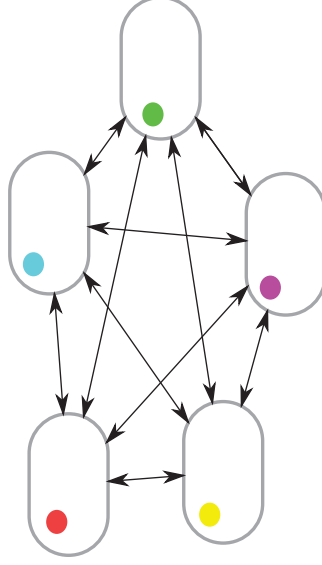
D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

MPI_Allgather

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala

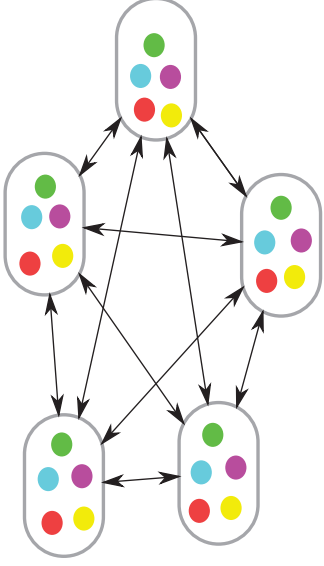




UPPSALA
UNIVERSITET

MPI_Allgather

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology



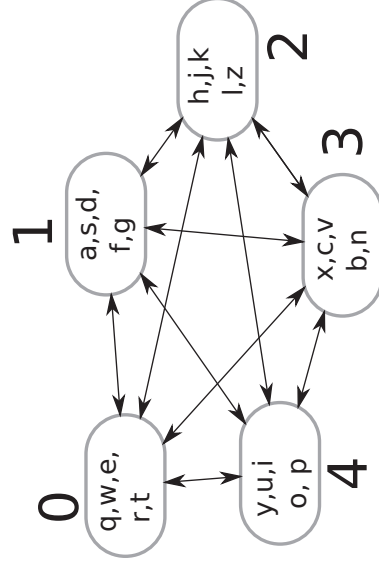
D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

MPI_Alltoall

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology



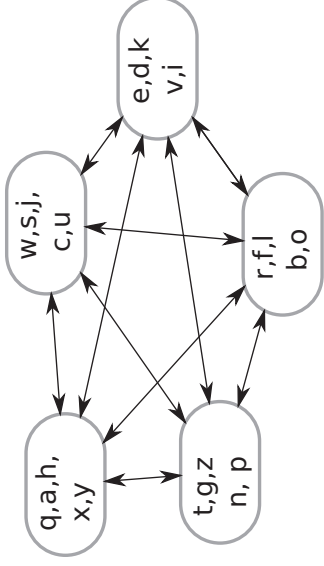
D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

MPI_Alltoall

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Timing in MPI

```
double start_time;  
start_time = MPI_Wtime();  
...  
// Compute  
...  
double finish_time;  
finish_time = MPI_Wtime();  
  
// Elapsed time  
double elapsed_time;  
elapsed_time = finish_time - start_time;
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

MPI Datatype

MPI_CHAR	signed char
MPI_SHORT	signed short int
MPI_INT	signed int
MPI_LONG	signed long int
MPI_UNSIGNED_CHAR	unsigned char
MPI_UNSIGNED_SHORT	unsigned short int
MPI_UNSIGNED	unsigned int
MPI_UNSIGNED_LONG	unsigned long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_LONG_DOUBLE	long double
MPI_PACKED	-
MPI_BYTE	-

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Derived Datatype

- ▶ Send non-contiguous data
- ▶ Reduce the number of communication calls
- ▶ `MPI_Type_vector()` – regular distribution
- ▶ `MPI_Type_indexed()` – irregular distribution
- ▶ `MPI_Type_struct()` – different datatypes

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes**
- Topology

- D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala


$$A = [1, 2, 3, 4; 5, 6, 7, 8; 9, 10, 11, 12];$$

- D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes**
- Topology

- ▶ `MPI_Datatype` `newtype`;
- ▶ `MPI_Type_vector(3,2,4,MPI_INT,&newtype)`;
- ▶ `MPI_Commit(&newtype)`;
- ▶ `MPI_Send(&A[0][1],1,newtype,1,0,comm)`;
- ▶ Sends new array [2 3 6 7 10 11] to process 1
- ▶ Could be interpreted as matrix with 3 rows, 2 cols



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes**
- Topology

- ▶ `MPI_Type_indexed (count, blocklens[], offsets[], old_type, newtype)`
- ▶ Make blocks at different offsets
- ▶ Displacements between successive blocks need not be equal



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

MPI_TYPE_INDEX

```
int numtasks, rank, source=0, dest, tag=1;
int blocklengths[2], displacements[2];
float a[16] =
{1.0, 2.0, 3.0, 4.0, 5.0, 6.0, 7.0, 8.0,
 9.0, 10.0, 11.0, 12.0, 13.0, 14.0, 15.0, 16.0};
float b[NELEMENTS];
MPI_Status stat;
MPI_Datatype indextype;

MPI_Init(&argc,&argv);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
MPI_Comm_size(MPI_COMM_WORLD, &numtasks);

blocklengths[0] = 4;
blocklengths[1] = 2;
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

MPI_TYPE_INDEX

```
displacements[0] = 5;
displacements[1] = 12;

MPI_Type_indexed(2, blocklengths, displacements,
                 MPI_FLOAT, &indextype);
MPI_Type_commit(&indextype);
if (rank == 0) {
    for (int i=0; i < numtasks; i++)
        MPI_Send(a, 1, indextype, i, tag, MPI_COMM_WORL
    }

MPI_Recv(b, NELEMENTS, MPI_FLOAT, source, tag,
         MPI_COMM_WORLD, &stat);
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

MPI_TYPE_INDEX

Output :

```
rank= 0  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 1  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 2  b= 6.0 7.0 8.0 9.0 13.0 14.0
rank= 3  b= 6.0 7.0 8.0 9.0 13.0 14.0
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline
Intro and motivation
Code Body
Communicators
Send and Receive
Other P2P Funcs
Global Funcs
Timing
Datatypes
Topology

Introduction and motivation

Code Body

Communicators

Send and Receive

Other Point-to-Point Functions

Global Functions

Timing

Datatypes

Topology

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala





UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

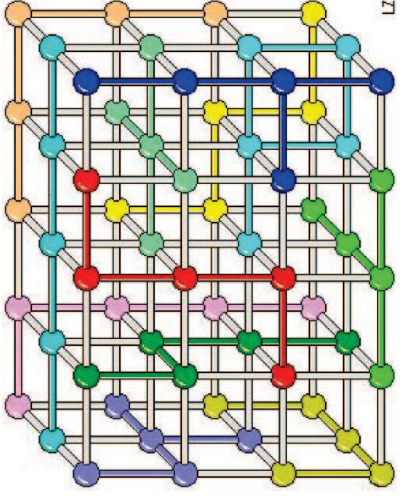
Global Funcs

Timing

Datatypes

Topology

Virtual Topology



D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UPPSALA
UNIVERSITET

Outline

Intro and motivation

Code Body

Communicators

Send and Receive

Other P2P Funcs

Global Funcs

Timing

Datatypes

Topology

Cartesian Topology

MPI_Cart_create(MPI_Comm old_comm, int ndims, int
*dim_size, int *periods, int reorder, MPI_Comm
*newcomm);

- ▶ Create a cartesian communicator
- ▶ with ndim (dimension) and dim_size (size)
- ▶ periods - 1 if some of the sides are periodic
- ▶ reorder - just ignore it

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



UNIVERSITÄT

- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology

```
old_comm = MPI_COMM_WORLD;
ndims = 2;
dim_size[0] = 3;
dim_size[1] = 2;
periods[0] = 1;
periods[1] = 0;
reorder = 1;
```

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



Map rank to coordinate

- ▶ MPI_Cart_coords();
- ▶ Input rank
- ▶ Output coordinate

- ▶ MPI_Cart_rank();
- ▶ Input coordinate
- ▶ Output rank

- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology**



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology**

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala



- Outline
- Intro and motivation
- Code Body
- Communicators
- Send and Receive
- Other P2P Funcs
- Global Funcs
- Timing
- Datatypes
- Topology**

D. Lukarski, J.Rantakokko, Jan, 2014, Uppsala

