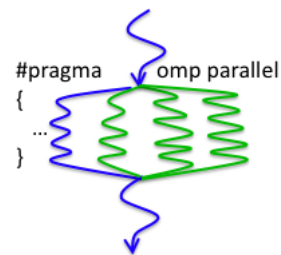


# Shared Memory (Multi-Core) Programming with OpenMP

Jarmo Rantakokko  
Senior lecturer, IT UU



**OpenMP:** Open specification for Multi Processing  
([www.openmp.org](http://www.openmp.org), v1.0 1997 - v4.0 2013, we will use v3.0)

Shared address space model, based on *threads*

**Thread:**

- Light weight process, global addresses
- Private program counter, independent
- Private stack pointer, private data

⇒ All threads have access to global data, can run in parallel and have some private data on stack.

On a multi-core node the threads are scheduled over the CPU's to the different cores.

⇒ One node on IT-servers can run 8 parallel threads.

Insert compiler directives for parallelization of Computations  $\Rightarrow$  high-level model

```
#pragma omp parallel for
for (i=2; i<=N-1; i++)
    A[i]=F(B[i-1]+B[i]+B[i+1]);
```

Loop is automatically parallelized over all threads, different iterations on different threads. Arrays A and B are global data, loop variable i is private.

```
NLOC=N/NPROC
ALLOCATE (A(NLOC),B(NLOC))
. . .
(Standard send/rcv avoiding deadlock)
```

```
IF (MOD(PID,2)==1) THEN
    CALL MPI_SEND(B(1),LEFT)
    CALL MPI_RECV(TEMP1,LEFT)
ELSEIF (MOD(PID,2)==0 AND PID<NPROC-1)
    CALL MPI_RECV(TEMP2,RIGHT)
    CALL MPI_SEND(B(NLOC),RIGHT)
END IF
IF (MOD(PID,2)==1 AND PID<NPROC-1) THEN
    CALL MPI_SEND(B(NLOC),RIGHT)
    CALL MPI_RECV(TEMP2,RIGHT)
ELSEIF (MOD(PID,2)==0 AND PID>0)
    CALL MPI_RECV(TEMP1,LEFT)
    CALL MPI_SEND(B(1),LEFT)
END IF
```

*(Simpler with non-blocking communication  
MPI\_Irecv followed by MPI\_Send)*

```
IF (PID>0) THEN
    A(1)=F(TEMP1+B(1)+B(2))
END IF
FOR (I=2; I<NLOC-1; I++)
    A(I)=F(B(I-1)+B(I)+B(I+1))
END DO
IF (PID<NPROC-1) THEN
    A(NLOC)=F(B(NLOC-1)+B(NLOC)+TEMP2)
END IF
```

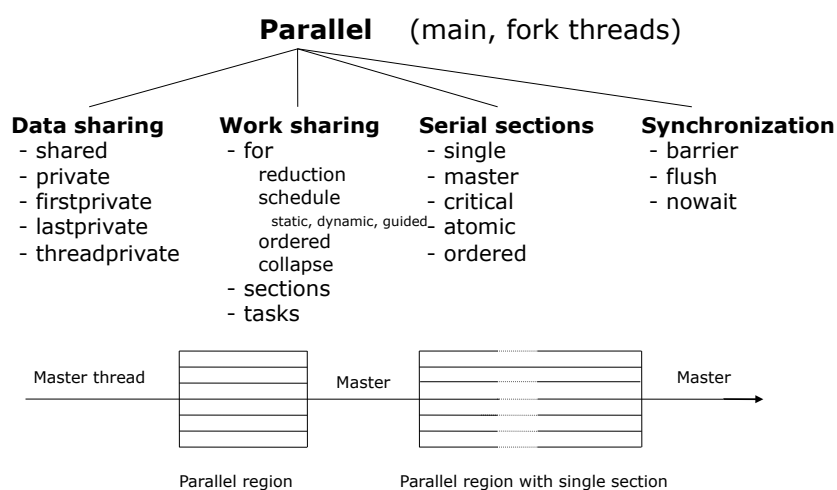




```
struct thread_data
{
    int j1;
    int j2;
};
void *compute(void *arg)
{
    int j1,j2;
    struct thread_data *index;
    index=(struct thread_data *)arg;
    j1=index->j1;
    j2=index->j2;

    for (j=j1;j<j2;j++)
        A[i]=F(B[i-1]+B[i]+B[i+1]);
}
int main(){
...
for(t=0; t<NUM_THREADS; t++) {
    index[t].j1=t*len/NUM_THREADS;
    index[t].j2=(t+1)*len/NUM_THREADS;
    pthread_create(&threads[t], &attr, compute,
        (void *) &index[t]); }
}
```

## OpenMP directives:



## OpenMP library functions:

- `omp_set_num_threads`
- `omp_get_num_threads`
- `omp_get_max_threads`
- `omp_get_thread_num`
- `omp_set_nested`
- and more (e.g. lock)

Allows for more flexible and user controlled (e.g. load balancing) programming than with the standard directives

### Environment variables: (export VARIABLE=value)

- `OMP_NUM_THREADS`
- `OMP_SCHEDULE`
- `OMP_NESTED`
- and more (stacksize, wait policy)

To run on 4 threads, before start of program do:  
`export OMP_NUM_THREADS=4`

## Directives: (Support only in Fortran/C/C++)

C/C++: `#pragma omp directive`  
`{ code block }`

Fortran: `!$omp directive`  
`code block`  
`!$omp end directive`

**Note:** The directives are ignored by non-supporting compiler or if OpenMP-flag is turned off in compiling.  
⇒ Portable code between single CPU, multi-core, and general parallel computers.

Also, possible to parallelize code incrementally  
(start with heaviest routine and continue until sufficient parallelism and performance are achieved)

## Parallel: (Fork-Join of threads)

```
#pragma omp parallel [subdirectives]
{
  "parallel code"
}
```

### Subdirectives:

```
if ( true/false )      -- parallel/serial
num_threads( int )    -- Number of threads
reduction (op:var)    -- parallel reduction
```

+ directives for data sharing (private/shared)

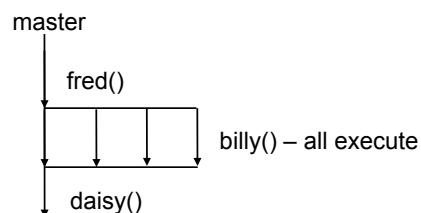
## Parallel: (Fork-Join of threads)

```
#pragma omp parallel
{
  "parallel code"
}
```

If no subdirectives, all data shared (global) and all code executed in parallel by all threads. At the end of parallel the threads are synchronized and joined.

Ex: program p1

```
...
call fred()
#pragma omp parallel
{ call billy() }
call daisy()
```



## Example HelloWorld:

```
#include <omp.h>
#include <stdio.h>

int main(int argc, char *argv[]) {

#pragma omp parallel
{
    printf("Hello world! %d\n",omp_get_thread_num());
}
}
```

**Task:** Compile and run the program helloworld.c

- gcc -fopenmp helloworld.c -o hello
- ./hello

Run on different number of threads.

In what ways can we change the number of threads?

## Data sharing:

- **shared**( [list of variables] ) - default
- **private**( [list of variables] )

Ex: program p2

```
...
a=100; b=0;
#pragma omp parallel private(a) num_threads(10)
{
    b=b+1000;
    a=b+a;
}
printf("a= %d, b= %d \n", a,b);
```

**Task:** What is the result, run the program datasharing.c several times and explain the output.

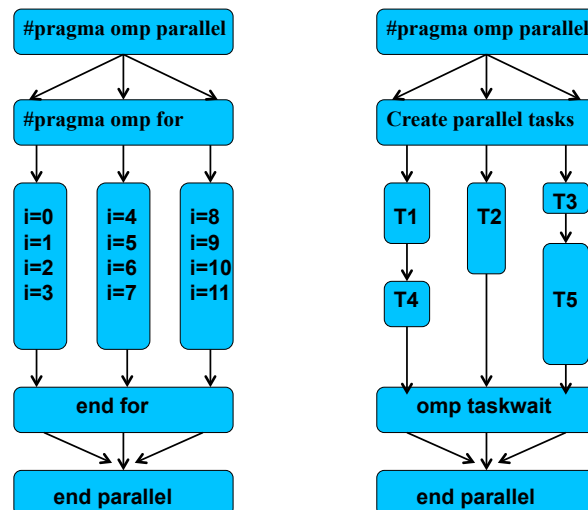
**Note:** All private variables are allocated on the stack  
=> uninitialized at entry and removed at exit,  
*original* a not equal to *private a*!

**Note2:** Shared variables must be protected from  
simultaneous writes by different threads!  
(Use critical directive or locks.)

- **firstprivate**( [list of variables] )  
As private but the variables are initialized from the original variable (in master) before parallel.
- **lastprivate**( [list of variables] )  
At exit, the original variable gets the value from the thread executing the last iteration in a loop using the for-directive or the last section in the sections-directive.
- **threadprivate**( [list of variables] )  
Make global file scope variables local and persistent to a thread through the execution of multiple parallel regions.

## Work sharing (within parallel)


- Loop level parallelism – **for**
- Task parallelism – **sections, tasks**



Institutionen för informationsteknologi | [www.it.uu.se](http://www.it.uu.se) | Jarmo Rantakokko

## for-directive:

```
#pragma omp for [subdirectives]
for (i=1; i<=n; i++)
{ loop-body }
```



### Subdirectives:

- Private
- Firstprivate
- Lastprivate
- Reduction
- Schedule
- Ordered
- Collapse

Without subdirectives, loop counter is private, loop space is divided statically into `nthr` equal pieces, and run in parallel (different iterations in different threads). Threads are synchronized at end of the for-directive.

**Note:** We must have a perfectly parallel loop!

Institutionen för informationsteknologi | [www.it.uu.se](http://www.it.uu.se) | Jarmo Rantakokko



## Example: Enumeration sort

```
for (j=0;j<len;j++)
{
    rank=0;
    for (i=0;i<len;i++)
        if (indata[i]<indata[j]) rank++;
    outdata[rank]=indata[j];
}
```

For each element (j) check how many other elements (i) are smaller than it => rank  
Perfectly parallel tasks for each element (j)

**Task:** Parallelize the j-loop in enumsort.c and set appropriate variables as private.

## Reduction( op:[list of variables] )

Performs a global reduction using  
**op**=+, -, \*, max, min, or a logical operator

```
sum=0;
#pragma omp parallel for reduction(+:sum)
for (i=0;i<n;i++) sum=sum+a[i];
```

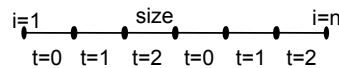
```
sum=0;
#pragma omp parallel private(locsum)
{
    locsum=0;
    #pragma omp for
    for (i=0;i<n;i++) locsum=locsum+a[i];
    #pragma omp critical
    { sum=sum+locsum; }
}
```

**Task:** Parallelize the inner i-loop in enumsort.c and compare the performance with your first parallelization.  
What are the performance obstacles and/or advantages?

## Schedule( type, [size] )

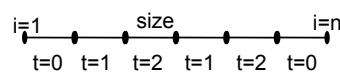
Divides the iteration space into chunks=size and schedules the chunks to threads according to type.  
(size=n/nthr by default)

### type=static:



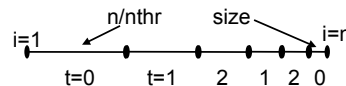
Assign the chunks cyclicly to threads

### type=dynamic:



Dynamic scheduling, as soon as a thread is ready it gets a new chunk

### type=guided:



As dynamic but the chunk size is decreasing towards end.  
Minimizes synchronization time.

### type=runtime:

Decide at runtime using the environment variable  
`export schedule=type` (where *type* is some above).

### type=auto:

Let the run-time system and/or compiler decide automatically.

**Note:** Static scheduling is good for data locality (cache) while dynamic/guided good for load balance.

**Task:** Try different scheduling options in the program `loop.c`, what gives the best performance, what is the theoretically minimal runtime, how can we get that?

## Ordered

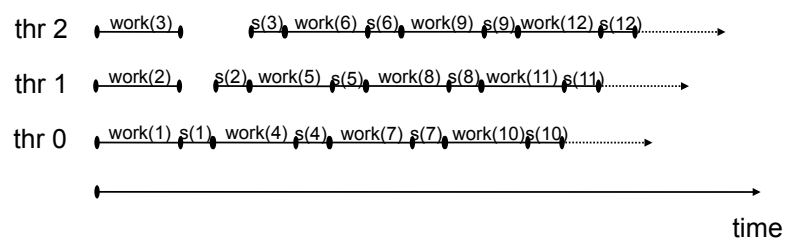
Only one thread is allowed to the ordered block at a time and sequentially in loop order. Useful for I/O.

```
#pragma omp parallel
{
    #pragma omp for schedule(static,1) ordered
    for (i=0;i<n;i++)
    {
        call work(i)          ! parallel work

        #pragma omp ordered
        { call s(i) }          ! serial section
    }
}
```

Assume  $\text{work}(i) \gg s(i) \Rightarrow$  Parallelism, pipelining effect

## Static,1:



What happens if we use default scheduling ( $\text{size} = n/\text{nthr}$ )?

## Collapse directive

Allow *collapsing* of perfectly nested loops, i.e., form a single loop and then parallelize that.

Example: parallelize both i and j-loop in MxM

```
#pragma omp parallel
{
  #pragma omp for collapse(2) private(i,j,k)
  for (i=0; i<len; i++)
    for (j=0; j<len; j++)
    {
      c[i*len+j]=0.0;
      for (k=0; k<len; k++)
        c[i*len+j]+=a[i*len+k]*b[k*len+j];
    }
}
```

## Task parallelism (static predefined tasks)

### Sections

<pre>#pragma omp sections [subdirectives] {   #pragma omp section   { task 1 }    #pragma omp section   { task 2 }    etc. }</pre>	<p>Subdirectives:</p> <ul style="list-style-type: none"> <li>• Private</li> <li>• Firstprivate</li> <li>• Lastprivate</li> <li>• Reduction</li> </ul>
--	---

The sections/tasks are scheduled (statically) to the threads and run in parallel. At end of sections the threads are synchronized. (No load balancing).

## Nested parallelism (load balancing of sections)

```
omp_set_nested(1);
#pragma omp parallel sections num_threads(2)
{
    #pragma omp section
    {
        #pragma omp parallel for num_threads(P1)
        for (k=0;k<n1;k++)
            call WORK1(A[K])
    }
    #pragma omp section
    {
        #pragma omp parallel for num_threads(P2)
        for (k=0;k<n2;k++)
            call WORK2(A[K])
    }
}
```

Assign appropriate number of threads to each section.

**Task:** Make a two-level parallelization of enumsort.c

## Task directive (dynamic tasks)

Can implement task-queues that are scheduled dynamically to all available threads in a parallel environment. Tasks can be generated at run-time.

Generate a task:

```
#pragma omp task [if/untied/'datasharing']
{ task }
```

Wait for all tasks to complete

```
#pragma omp taskwait
```

Task scheduling points at following locations:

1. Generation of task
2. Last instruction in task
3. Taskwait-directive
4. Implicit and explicit barriers

**Task:** Study and run the program task.c. What is the effect of if and nowait?

## Serial sections

Avoid terminating threads, lose data in cache if threads rescheduled to different CPUs or cores (with fork-join model).

### **#pragma omp single [subdirectives]**

The code-block within single is executed only by one thread, the others skip and wait at the end of block.

Subdirectives: - private  
                  - firstprivate

### **#pragma omp master**

The code-block is executed only by master thread, the other skip and continue (no barrier).

### **#pragma omp critical [name]**

The code-block is executed by one thread at a time.

As ordered but no predefined order.

If no name all critical sections have the same name.

Only one critical section with the same name can be executed by one thread at a time.

### **#pragma omp atomic**

Atomic update by one thread at a time. As critical but applies only for a one line expression. (Does not include a memory flush.)

## Synchronization

Done implicitly at end of:

- parallel
- for
- sections
- single

Explicit barrier:

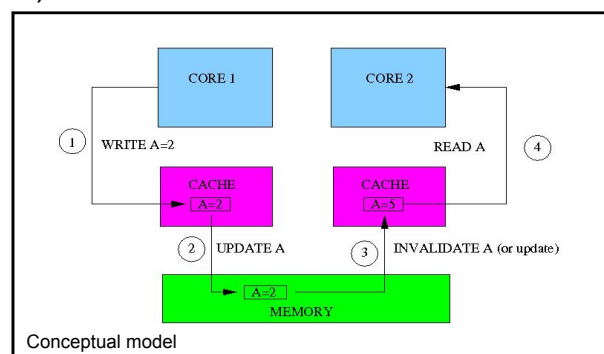
`#pragma omp barrier`

Can override with *nowait*:

```
#pragma omp for nowait
for (i=0;i<n;i++)
{ code }
```

**Note:** If *nowait* be careful not to use data updated by other threads, *nowait* overrides memory flush!

In a **memory flush** all thread visible shared variables are refreshed (caches invalidated and memory updated). A memory flush is performed at all barriers (implicit and explicit) and before/after a critical directive.



=> Be very careful with *nowait* !!! (*Nowait* removes 2 & 3)

OpenMP has a *relaxed-consistency* model, i.e., the threads can cache data not keeping exact consistency.

**Task:** Run the program `memory.c` and explain its output. How can we fix the problem?

```
int main(int argc, char *argv[]) {
    int id,nthr;

    #pragma omp parallel private(id)
    {
        nthr=-1;
        id=omp_get_thread_num();

        #pragma omp single
        { nthr=omp_get_num_threads(); }

        printf("Hello world! %d %d\n",id,nthr);
    }
}
```

## Synchronization with locks

Can lock a code section and/or data only accessible to a specific thread. Routines include a flush.

`omp_init_lock` - `omp_destroy_lock`  
`omp_set_lock` - `omp_unset_lock`  
`omp_test_lock`

Example:

```
omp_lock_t lockvar;
omp_init_lock(lockvar);
...
#pragma omp parallel {
    ...
    omp_set_lock(lockvar);
    sum=sum+a;
    omp_unset_lock(lockvar);
}
```