

Case Study: Implementing Enumeration Sort in OpenMP

Jarmo Rantakokko
Senior lecturer, IT UU



Purpose: To study and identify different parallel overheads in OpenMP (we are not interested in how to parallelize enumsort in the best way)

Algorithm: Enumeration Sort

```
for (j=0;j<len;j++)  
{  
    rank=0;  
    for (i=0;i<len;i++)  
        if (indata[i]<indata[j]) rank++;  
    outdata[rank]=indata[j];  
}
```

For each element (j) check how many other elements (i) are smaller than it => rank
Perfectly parallel tasks for each element (j)

Alternative 1: Parallelize j-loop

```
#pragma omp parallel for private(rank,i)
for (j=0;j<len;j++)
{
    rank=0;
    for (i=0;i<len;i++)
        if (indata[i]<indata[j]) rank++;
    outdata[rank]=indata[j];
}
```

Note 1: If rank equal on two threads
⇒ Race condition (but write same data)

Note 2: All threads reading all data for each element
⇒ Mem BW limited performance (if data does not fit in cache, especially bad on NUMA)

Alternative 2: Parallelize i-loop

```
for (j=0;j<len;j++)
{
    rank=0;
    #pragma omp parallel for reduction (+:rank)
    for (i=0;i<len;i++)
        if (indata[i]<indata[j]) rank++;
    outdata[rank]=indata[j];
}
```

Note 1: Frequent creation/termination of threads and small tasks per thread (high parallel overhead)

Note 2: Each thread works only on a part of the data in all iterations, good for cache performance (if the whole array does not fit in cache).
Also no race condition, only master updates

Alternative 3: Use one parallel region

```
#pragma omp parallel private(j)
{
  for (j=0;j<len;j++)
  {
    #pragma omp single
    { rank=0; }

    #pragma omp for reduction (+:rank)
    for (i=0;i<len;i++)
      if (indata[i]<indata[j]) rank++;

    #pragma omp single
    { outdata[rank]=indata[j]; }
  }
}
```

Note: 3 barriers per iteration, how can we decrease the number of synchronization points?

Alternative 4: Interleave two iterations

```
j1=0; j2=1; rank1=0; rank2=0;
#pragma omp parallel
{
  while (j1<len)
  {
    #pragma omp for reduction (+:rank1) // Barrier
    for (i=0;i<len;i++)
      if (indata[i]<indata[j1]) rank1++;

    #pragma omp single nowait
    { outdata[rank1]=indata[j1];
      rank1=0; j1+=2; }

    if (j2>=len) break;

    #pragma omp for reduction (+:rank2) // Barrier
    for (i=0;i<len;i++)
      if (indata[i]<indata[j2]) rank2++;

    #pragma omp single nowait
    { outdata[rank2]=indata[j2];
      rank2=0; j2+=2; }
  }
}
```

Results (runtime):

Nthr	Enum1	Enum2	Enum3	Enum4
1	11.5	13.0	11.7	12.0
2	5.80	7.90	9.23	7.80
4	2.96	5.73	9.05	5.57
8	1.50	7.90	38.1	6.13
16	0.75	NA	NA	NA

What overheads do we have?


Enum1: All threads read all data in all iterations, mem BW limited performance (if small cache).

Enum2: Create/terminate threads in each iteration

Enum3: Three barriers per iteration

Enum 4: One barrier per iteration => Memory flush
We have update of invalid cache-lines, outdata is updated irregularly (randomly) and we get "communication" due to (false) sharing in outdata

Alternative 5: Let each thread be responsible for a fixed section of outdata and only that thread writes in the corresponding locations



Alternative 5: Owner writes

```

j1=0; j2=1; rank1=0; rank2=0;
#pragma omp parallel
{
  while (j1<len)
  {
    #pragma omp for reduction (+:rank1)
    for (i=0;i<len;i++)
      if (indata[i]<indata[j1]) rank1++;

    if (rank1/(len/nthr)==thrid)
    { outdata[rank1]=indata[j1];
      rank1=0; j1+=2; }

    if (j2>=len) break;

    #pragma omp for reduction (+:rank2)
    for (i=0;i<len;i++)
      if (indata[i]<indata[j2]) rank2++;

    if (rank2/(len/nthr)==thrid)
    { outdata[rank2]=indata[j2];
      rank2=0; j2+=2; }
  }
}

```

Results (runtime):

Nthr	Enum1	Enum2	Enum3	Enum4	Enum5
1	11.5	13.0	11.7	12.0	11.4
2	5.80	7.90	9.23	7.80	6.90
4	2.96	5.73	9.05	5.57	5.20
8	1.50	7.90	38.1	6.13	5.75
16	0.75	NA	NA	NA	NA

Enum1: All data fits in cache, perfectly parallel
Maximal work per thread, minimal parallel overhead

Enum5: Still one barrier per iteration and small
work load per thread between synchronizations
(we need to go up in problem size to pay off)

Alternative 6: Nested parallelism

```
omp_set_nested(1);  
#pragma omp parallel for private(rank) num_threads(4)  
for (j=0;j<len;j++)  
{  
    rank=0;  
    #pragma omp parallel for reduction (+:rank) num_threads(2)  
    for (i=0;i<len;i++)  
        if (indata[i]<indata[j]) rank++;  
    outdata[rank]=indata[j];  
}
```

Note: Increase the parallel overhead compared to Enum1
(create/terminate threads in each iteration j)
Decrease the parallel overhead compared to Enum2
(synchronize a smaller team of threads)